

HW1 Report for CSCI544 SPRING23

After the Dataset Preparation Step, I formed a balanced dataset by merging the 3 classes. Proceeding from this point, I used some conventional data cleaning and preprocessing techniques to clear the noises in the dataset, thus leading to better performance of my trained classification models. Things I did are listed below in order.

1. Removed unwanted Html tags using BeautifulSoup, as they are not involved in sentiment analysis.
2. Added spell correction to correct mistyped English texts. Reviewers often mistype words unintentionally, so they shouldn't be treated as different words or discarded. It's logical to correct the mistyped words to the reviewer's intended format, as this is for sentiment analysis.
3. Extended contracted words for easier preprocessing.
4. Removed non-alphabetical characters as we're only interested in English Review.
5. Converted to lower cases. Words like Book and book mean the same but when not converted to lowercase those two are represented as two different words.
6. Removed stop words to give more focus to the important information. We're using TF-IDF so removing stopwords is a good idea.
7. Converted words to their lemmatized format. Grouping different inflected forms of words into the root form, having the same meaning.
8. Removed extra whitespaces.

Other things I did: Added text translation to translate non-English text to English. Not included as it either requires an internet connection, or the library is tremendously large.

Training Classification Models:

Extracted TF-IDF features and split them into a training set and testing set. Trained 4 models using Ski-learn library.

Parameters setting:

Perceptron: Default

Linear SVM: Using SVM's conventional loss function "hinge"

Logistic Regression: Using "sag" solver for fast convergence

Multinomial Naïve Bayes: Set fit_prior to False as it gives better result

TF-IDF: Using 5-grams as it gives better result

The training results are shown in the below section.

Import Libraries

```
In [1]: # Python Version: 3.7.9
# Pandas Version: 1.3.5
# BeautifulSoup Version: 4.11.1
# Contractions Version: 0.0.18
# Setuptools Version: 60.2.0
# SymSpellpy Version: 6.7.7
# NLTK Version: 3.8.1
# Scikit-Learn Version: 1.0.2

import pandas as pd
import re
from bs4 import BeautifulSoup
import contractions as ct
import pkg_resources
from symspellpy import SymSpell
import warnings

from nltk.corpus import wordnet as wn
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from nltk import map_tag, WordNetLemmatizer, pos_tag

from sklearn.svm import LinearSVC
from sklearn.metrics import classification_report
from sklearn.naive_bayes import MultinomialNB
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import Perceptron, LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
```

Define Functions

```
In [2]: # Drop empty & duplicated rows
def init_data(data_frame):
    data_frame.dropna(inplace=True)
    data_frame.drop_duplicates(inplace=True)
    data_frame['star_rating'] = data_frame['star_rating'].astype('int')
    return data_frame

# Init spell checker object
def init_spell_checker():
    sym_spell_obj = SymSpell(max_dictionary_edit_distance=2, prefix_length=7)
    dictionary_path = pkg_resources.resource_filename(
        "symspellpy", "frequency_dictionary_en_82_765.txt"
    )
    bigram_path = pkg_resources.resource_filename(
        "symspellpy", "frequency_bigramdictionary_en_243_342.txt"
    )
    sym_spell_obj.load_dictionary(dictionary_path, term_index=0, count_index=1)
    sym_spell_obj.load_bigram_dictionary(bigram_path, term_index=0, count_index=2)

    return sym_spell_obj

# Spell correct the input text
def spell_correct(text):
    input_term = text
    suggestions = sym_spell.lookup_compound(
        input_term, max_edit_distance=2, transfer_casing=True
    )
    return suggestions[0].term

# Lemmatize the input word
def word_lemmatization(word):
    treebank_pos_tag = pos_tag([word])[0][1]
    universal_pos_tag = map_tag('en-ptb', 'universal', treebank_pos_tag)

    if universal_pos_tag == "ADJ":
        word = wn1.lemmatize(word, wn.ADJ)
    elif universal_pos_tag == "VERB":
        word = wn1.lemmatize(word, wn.VERB)
    elif universal_pos_tag == "NOUN":
        word = wn1.lemmatize(word, wn.NOUN)
    elif universal_pos_tag == "ADV":
        word = wn1.lemmatize(word, wn.ADV)
    word = get_adverb_lemma(word)

    return word

# Get an input adverb's Lemma if any
def get_adverb_lemma(word):
    has_suggestion = False
    param_wn_synset = word + ".r.01"

    # check if word's synset contains adverb option
    for i in range(0, len(wn.synsets(word))):
        if param_wn_synset == str((wn.synsets(word)[i])).split("\'")[1]:
            has_suggestion = True

    if not has_suggestion:
```

```

        return word

# if yes and suggestion not empty then return suggestion, else return original word
suggest_lemma_list = wn.synset(param_wn_synset).lemmas()[0].pertainyms()
if len(suggest_lemma_list) > 0:
    return suggest_lemma_list[0].name()
else:
    return word

# Remove stop words and Lemmatize the remaining words
def lemmatize_non_stopwords(review_body_string):
    word_tokens = word_tokenize(review_body_string)
    buffer_string = ""

    for w in word_tokens:
        if w not in stop_words:
            w = word_lemmatization(w)
            buffer_string = buffer_string + w + " "

    buffer_string = re.sub(' +', ' ', buffer_string).strip()
    return buffer_string

# Data cleaning & preprocessing
def data_cleaning(data_frame):
    before_data_cleaning_reviews_total_length = 0
    after_data_cleaning_reviews_total_length = 0
    before_data_preprocessing_reviews_total_length = 0
    after_data_preprocessing_reviews_total_length = 0

    for i in range(0, len(data_frame)):

        if data_frame['star_rating'][i] == '1' or data_frame['star_rating'][i] == '2':
            data_frame.loc[i, ['star_rating']] = 'Class 1'
        elif data_frame['star_rating'][i] == '3':
            data_frame.loc[i, ['star_rating']] = 'Class 2'
        elif data_frame['star_rating'][i] == '4' or data_frame['star_rating'][i] == '5':
            data_frame.loc[i, ['star_rating']] = 'Class 3'

        review_text = data_frame['review_body'][i]
        before_data_cleaning_reviews_total_length = before_data_cleaning_reviews_total_length + len(review_text)

        # remove un-wanted html tags
        if BeautifulSoup(review_text, "html.parser").find():
            review_text = BeautifulSoup(review_text, "html.parser").get_text(" ")

        # spell correction
        review_text = spell_correct(review_text)

        # text extend contractions
        review_text = ct.fix(review_text)

        # remove non-alphabetical chars
        regex = re.compile('[^a-zA-Z]')
        review_text = regex.sub(' ', review_text)

        # convert to lower case
        review_text = review_text.lower().strip()
        review_text = " ".join(review_text.split())

        # end of data cleaning, before data processing
        after_data_cleaning_reviews_total_length = after_data_cleaning_reviews_total_length + len(review_text)

        # start of data processing
        before_data_preprocessing_reviews_total_length = before_data_preprocessing_reviews_total_length + len(
            review_text)
        review_text = lemmatize_non_stopwords(review_text)
        # end of data processing
        review_text = " ".join(review_text.split())
        after_data_preprocessing_reviews_total_length = after_data_preprocessing_reviews_total_length + len(review_text)

        data_frame.loc[i, ['review_body']] = review_text

    print("Average length of reviews before data cleaning: " + str(before_data_cleaning_reviews_total_length / len(
        data_frame)) + ", Average length of reviews after data cleaning: " + str(
        after_data_cleaning_reviews_total_length / len(data_frame)))
    print("Average length of reviews before data preprocessing: " + str(
        before_data_preprocessing_reviews_total_length / len(
            data_frame)) + ", Average length of reviews after data preprocessing: " + str(
            after_data_preprocessing_reviews_total_length / len(data_frame)))
    print("\n")

    return data_frame

# Print the training result
def generate_report(y_test, y_pred):
    report = classification_report(y_test, y_pred, zero_division=1, output_dict=True)
    print("Class 1 Precision: " + str(report['Class 1']['precision']) + ", Class 1 Recall: " + str(
        report['Class 1']['recall']) + ", Class 1 f1-score: " + str(report['Class 1']['f1-score']))
    print("Class 2 Precision: " + str(report['Class 2']['precision']) + ", Class 2 Recall: " + str(
        report['Class 2']['recall']) + ", Class 2 f1-score: " + str(report['Class 2']['f1-score']))
    print("Class 3 Precision: " + str(report['Class 3']['precision']) + ", Class 3 Recall: " + str(
        report['Class 3']['recall']) + ", Class 3 f1-score: " + str(report['Class 3']['f1-score']))
    print("Average Precision: " + str(report['macro avg']['precision']) + ", Averagage Recall: " + str(
        report['macro avg']['recall']) + ", Averagage f1-score: " + str(
            report['macro avg']['f1-score']))
    print("\n")

```

Initialization

```
In [3]: # Init
RANDOM_SAMPLE_SIZE = 20000
warnings.filterwarnings("ignore", category=UserWarning, module='bs4')
sym_spell = init_spell_checker()
stop_words = set(stopwords.words('english'))
wnl = WordNetLemmatizer()
```

Read Data

```
In [4]: # Reading data from cache. (data.pkl was generated by reading the given Amazon's dataset provided in HW1 description)
df = pd.read_pickle("./data.pkl")
df = init_data(df).reset_index(drop=True)
```

Form three classes and select 20000 reviews randomly from each class.

```
In [5]: # 3-classes and concat into 1
class1_df = df[df['star_rating'] <= 2].sample(RANDOM_SAMPLE_SIZE)
class2_df = df[df['star_rating'] == 3].sample(RANDOM_SAMPLE_SIZE)
class3_df = df[df['star_rating'] >= 4].sample(RANDOM_SAMPLE_SIZE)

balanced_df = pd.concat([class1_df, class2_df, class3_df]).reset_index(drop=True)
balanced_df['star_rating'] = balanced_df['star_rating'].astype('string')
```

Data Cleaning & Pre-processing

1. Remove un-wanted Html tags
2. Spell corrections
3. Text contractions
4. Remove non-alphabetical chars
5. Convert to lower cases
6. Remove stop words
7. Lemmatisation

```
In [6]: cleaned_balanced_df = data_cleaning(balanced_df)
# cleaned_balanced_df cache
cleaned_balanced_df.to_pickle('cleaned_balanced_df_official.pkl')
cleaned_balanced_df = pd.read_pickle("./cleaned_balanced_df_official.pkl")
```

Average length of reviews before data cleaning: 278.5658833333333, Average length of reviews after data cleaning: 269.23761666666667
Average length of reviews before data preprocessing: 269.23761666666667, Average length of reviews after data preprocessing: 154.18928333333332

TF-IDF Feature Extraction

```
In [7]: # tf-idf feacture matrix
tf_idf = TfidfVectorizer(lowercase=False, ngram_range=(1, 5))
tf_idf_result = tf_idf.fit_transform(cleaned_balanced_df['review_body'])
```

Split dataset into Training and Testing Set

```
In [8]: X_train, X_test, y_train, y_test = train_test_split(tf_idf_result, cleaned_balanced_df['star_rating'], test_size=0.2)
```

Perceptron

```
In [9]: # Train Perceptron Model & generate training report
clf_perceptron = Perceptron()
clf_perceptron = clf_perceptron.fit(X_train, y_train)
y_pred_perceptron = clf_perceptron.predict(X_test)
generate_report(y_test, y_pred_perceptron)
```

Class 1 Precision: 0.6571160169093471, Class 1 Recall: 0.6856162705219309, Class 1 f1-score: 0.671063676699844
Class 2 Precision: 0.5603917301414582, Class 2 Recall: 0.5175879396984925, Class 2 f1-score: 0.5381400208986415
Class 3 Precision: 0.707083128381702, Class 3 Recall: 0.7298806803757298, Class 3 f1-score: 0.7183010618363522
Average Precision: 0.6415302918108358, Averagage Recall: 0.6443616301987177, Averagage f1-score: 0.6425015864782794

SVM

```
In [10]: # Train SVM Linear Model & generate training report
clf_linear_svc = LinearSVC(loss='hinge')
clf_linear_svc = clf_linear_svc.fit(X_train, y_train)
y_pred_linear_svc = clf_linear_svc.predict(X_test)
generate_report(y_test, y_pred_linear_svc)
```

Class 1 Precision: 0.7024064808196331, Class 1 Recall: 0.7223719676549866, Class 1 f1-score: 0.7122493355883065
Class 2 Precision: 0.6075845012366035, Class 2 Recall: 0.5555276381909547, Class 2 f1-score: 0.5803911274445465
Class 3 Precision: 0.7346301633045149, Class 3 Recall: 0.7765930439197766, Class 3 f1-score: 0.755029001604344
Average Precision: 0.6815403817869171, Averagage Recall: 0.6848308832552393, Averagage f1-score: 0.682556488212399

Logistic Regression

```
In [11]: # Train Logistic Regression Model & generate training report
clf_logistic_regression = LogisticRegression(solver='sag')
clf_logistic_regression = clf_logistic_regression.fit(X_train, y_train)
y_pred_logistic_regression = clf_logistic_regression.predict(X_test)
generate_report(y_test, y_pred_logistic_regression)
```

Class 1 Precision: 0.700441609421001, Class 1 Recall: 0.699583435432492, Class 1 f1-score: 0.7000122594090965
Class 2 Precision: 0.5834348355663824, Class 2 Recall: 0.6017587939698492, Class 2 f1-score: 0.5924551638837353
Class 3 Precision: 0.7512437810945274, Class 3 Recall: 0.7283574511297284, Class 3 f1-score: 0.7396236143335911
Average Precision: 0.6783734086939702, Averagage Recall: 0.6765665601773566, Averagage f1-score: 0.6773636792088076

Naive Bayes

```
In [12]: # Train MultinomialNB Model & generate training report
clf_multinomial_nb = MultinomialNB(fit_prior=False)
clf_multinomial_nb = clf_multinomial_nb.fit(X_train, y_train)
y_pred_multinomial_nb = clf_multinomial_nb.predict(X_test)
generate_report(y_test, y_pred_multinomial_nb)
```

Class 1 Precision: 0.7267833109017496, Class 1 Recall: 0.6616025483950012, Class 1 f1-score: 0.6926629040533607
Class 2 Precision: 0.573621103117506, Class 2 Recall: 0.6010050251256281, Class 2 f1-score: 0.5869938650306749
Class 3 Precision: 0.7297691373025517, Class 3 Recall: 0.7623762376237624, Class 3 f1-score: 0.7457164142041223
Average Precision: 0.6767245171072691, Averagage Recall: 0.6749946037147971, Averagage f1-score: 0.675124394429386

Authorship

Author: Che Wei Wu Date: Jan 24, 2023 Description: The source code for USC_CSCI544_SPRING23_HW1

 JW